

# OS Project 2

## System Call on xv6 and User Defined Shell



**Prof. Yongtae Kim**

Computer Science and Engineering  
Kyungpook National University

# What is xv6 Operating System?

## ■ xv6 was developed at MIT for OS education

- It is run on multiprocessor Intel x86 and RISC-V systems
- It contains important Unix concepts and constructs
- It is an open-source project and can be compiled using the GNU C compiler
- It is normally run using the QEMU emulator, which is a free and open-source hardware emulator



## ■ Before you start (important!)

- xv6 runs on only x86 or RISC-V processor architecture
- You must use a x86-based computer, **not on apple silicon (M1 or M2) Mac**
- You are recommended to use VirtualBox with fresh installed Ubuntu 22.04
- You can use the VirtualBox used in OS Project 1 (new kernel equipped Ubuntu) and may see some glitches

# Install QEMU and build xv6

## ■ Install packages QEMU

- `$ sudo apt-get update`
- `$ sudo apt-get install git build-essential qemu-kvm gdb vim -y`

## ■ Download and build xv6

- `$ git clone https://github.com/mit-pdos/xv6-public.git`
- `$ cd xv6-public`
- `$ make qemu` // build and run xv6 on QEMU command
- You will see a new pop-up window for QEMU and you can type a xv6 command and see the results on either the original Linux terminal or QEMU terminal

```

.o sleeplock.o spinlock.o string.o switch.o syscall.o sysfile.o sysproc.o trapasm
.o trap.o uart.o vectors.o vm.o -b binary initcode entryother
objdump -S kernel > kernel.asm
objdump -t kernel | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$/d' > kernel.sym
dd if=/dev/zero of=xv6.img count=10000
10000+0 records in
10000+0 records out
5120000 bytes (5.1 MB, 4.9 MiB) copied, 0.0350795 s, 146 MB/s
dd if=bootblock of=xv6.img conv=notrunc
1+0 records in
1+0 records out
512 bytes copied, 0.000134223 s, 3.8 MB/s
dd if=kernel of=xv6.img seek=1 conv=notrunc
393+1 records in
393+1 records out
201436 bytes (201 kB, 197 KiB) copied, 0.000903657 s, 223 MB/s
qemu-system-i386 -serial mon:stdio -drive file=fs.img,index=1,media=disk,format=
raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -m 512
xv6...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap star
t 58
init: starting sh
$
  
```

Linux  
Terminal

```

Machine View
SeaBIOS (version 1.15.0-1)

iPXE (https://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8B590+1FECB590 CA00

Booting from Hard Disk...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap star
t 58
init: starting sh
$
  
```

QEMU  
Terminal

# System Call (1)

## ■ Making a new system call

- New system call is to get the current UTC time and return it to the user program
- `$ grep -n uptime *. [chS]`

```
mj@mj:~/Desktop/new/xv6-public$ grep -n uptime *. [chS]
syscall.c:105:extern int sys_uptime(void);
syscall.c:121:[SYS_uptime] sys_uptime,
syscall.h:15:#define SYS_uptime 14
sysproc.c:83:sys_uptime(void)
user.h:25:int uptime(void);
usys.S:31:SYSCALL(uptime)
mj@mj:~/Desktop/new/xv6-public$
```

- You will need to modify `syscall.c`, `syscall.h`, `sysproc.c`, `user.h`, `usys.S` for new system call
- In addition, you will need to write `date.c` for the new system call testing

# System Call (2)

## ■ Making date system call

- Writing date system call code in sysproc.c

```
int
sys_date(void)
{
    struct rtcdate* r;

    argptr(0, (void *)&r, sizeof(r));
    cmostime(r);

    return 0;
}
```

- Adding the date system call in syscall.c and syscall.h

extern int sys_write(void);	[SYS_mkdir] sys_mkdir,	#define SYS_mkdir 20
extern int sys_uptime(void);	[SYS_close] sys_close,	#define SYS_close 21
extern int sys_date(void);	[SYS_date] sys_date,	#define SYS_date 22

# System Call (3)

## ■ Preparing the date system call testing

### — Writing date.c

```
#include "types.h"
#include "user.h"
#include "date.h"

int main(int argc, char *argv[])
{
    struct rtcdate r;

    if (date(&r))
        printf(2, "date failed\n");
    else
        printf(1, "UTC: %d-%d-%dT%d:%d:%d+09:00\n",
            r.year, r.month, r.day, r.hour, r.minute, r.second);
    exit();
}
```

### — Adding the date system call interface to user.h and usys.S

```
int sleep(int);
int uptime(void);
int date(struct rtcdate *);
```

```
SYSCALL(sleep)
SYSCALL(uptime)
SYSCALL(date)
```

# System Call (4)

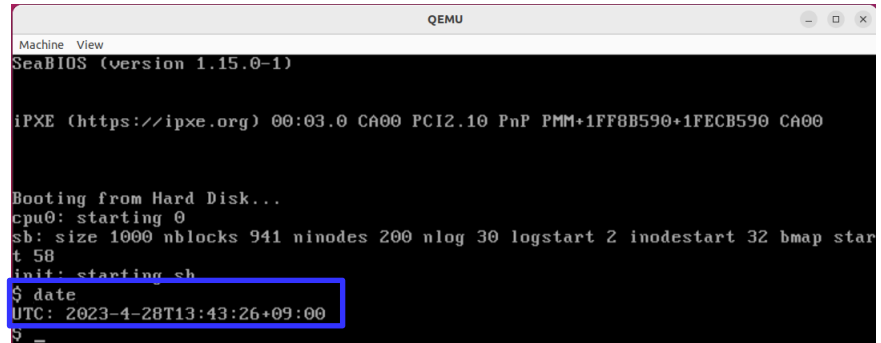
## ■ Compiling xv6 with new system call

- Modifying Makefile

```

UPROGS=\
    _cat\
    _echo\
    _forktest\
    _grep\
    _init\
    _kill\
    _ln\
    _ls\
    _mkdir\
    _rm\
    _sh\
    _stressfs\
    _usertests\
    _wc\
    _zombie\
    _date\

```



```

Machine View
SeaBIOS (version 1.15.0-1)

iPXE (https://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8B590+1FECB590 CA00

Booting from Hard Disk...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap star
t 58
init: starting sh
$ date
UTC: 2023-4-28T13:43:26+09:00
$ _

```

## ■ Build xv6 and test

- \$ make qemu
- Type \$date on qemu terminal (or Linux terminal) and see the time

## ■ Task #1: Make your own system call on xv6


# System Call Tracing

- **Modifying the xv6 kernel to trace each system call invocation**
  - Printing out the system call name & no, process id & name of each invocation
  - To do this, syscall.c should be modified

## Initial code

```
void
syscall(void)
{
    int num;
    struct proc *curproc = myproc();

    num = curproc->tf->eax;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        curproc->tf->eax = syscalls[num]();
    } else {
        cprintf("%d %s: unknown sys call %d\n",
            curproc->pid, curproc->name, num);
        curproc->tf->eax = -1;
    }
}
```



```
void
syscall(void)
{
    int num;
    struct proc *curproc = myproc();

    num = curproc->tf->eax;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        curproc->tf->eax = syscalls[num]();
        switch(num)
        {
            // your code here

            case 22: cprintf("date->pid: %d, name: %s, syscallno: %d\n",
                curproc->pid, curproc->name, num);
                break;
        }
    } else {
        cprintf("%d %s: unknown sys call %d\n",
            curproc->pid, curproc->name, num);
        curproc->tf->eax = -1;
    }
}
```

- `$ make qemu`

- **Task #2: Make it trace all of the xv6's system calls**



# User Defined Shell on Linux (1)

## ■ Download the example shell code

- Note that this shell runs on Linux, not on xv6
- `$ gcc sh.c // you will see a warning, which can be ignored`
- `$ ./a.out`

```
mj@mj:~/Desktop/new/sh$ ./a.out
COMP0312$ ls
exec not implemented
COMP0312$ mkdir
exec not implemented
COMP0312$
```

## ■ Take #3-1: Make the shell to execute commands by filling codes

- Hint: use `execvp` system call to implement the command execution

```
case ' ':
    ecmd = (struct execcmd*)cmd;
    if(ecmd->argv[0] == 0)
        exit(0);
    fprintf(stderr, "exec not implemented\n");
    // Remove above line and your code here ...
    break;
```



```
case ' ':
    ecmd = (struct execcmd*)cmd;
    if(ecmd->argv[0] == 0)
        exit(0);
    

Your code here


    break;
```

# User Defined Shell on Linux (2)

## I/O direction and pipe

- < redirection and pipe (|) were already implemented but > redirection wasn't
- \$ ./a.out // to do this, Task #3-1 must be implemented

```
COMP0312$ cat < sh.c | sort
```

```
:
```

```
void
```

```
    wait(&r);
```

```
    while(getcmd(buf, sizeof(buf)) >= 0){
```

```
        while(!peek(ps, es, "|")){
```

```
            while(peek(ps, es, "<>")){
```

```
                while(s < es && strchr(whitespace, *s))
```

```
                while(s < es && strchr(whitespace, *s))
```

```
                while(s < es && strchr(whitespace, *s))
```

```
                while(s < es && !strchr(whitespace, *s) && !strchr(symbols, *s))
```

```
COMP0312$ cat < sh.c | sort > test
```

```
> redir not implemented
```

```
COMP0312$
```

```
case '>':
    fprintf(stderr, "> redir not implemented\n");
    // Remove above line and your code here ...
    break;
```

```
case '<':
    rcmd = (struct redircmd*)cmd;
    if( (rcmd->fd = open(rcmd->file, rcmd->mode, 0644)) <= 0 )
    {
        fprintf(stderr, "file open error\n");
        close(rcmd->fd);
        break;
    }
    freopen(rcmd->file, "r", stdin);
    runcmd(rcmd->cmd);
    break;
```

```
case '|':
    pcmd = (struct pipecmd*)cmd;
    if( pipe(p) == -1 )
    {
        fprintf(stderr, "pipe error\n");
        break;
    }
    r = fork1();
    if( r == 0 )
    {
        dup2(p[1], 1);
        close(p[0]);
        runcmd(pcmd->left);
    }
    else
    {
        dup2(p[0], 0);
        close(p[1]);
        runcmd(pcmd->right);
    }
}
```

## Task #3-2: Make > redirection possible

- Hint: < redirection code

# OS Project 2

## ■ What to Do

- **Task #1:** Write your own system call on xv6
  - Your system call (proj2call) simply prints the following message to kernel:  
**COMP0312\_OS\_PROJ2\_yourStudentID\_yourName: Hello xv6**
- **Task #2:** Write codes to trace all xv6's system call after booting
  - Tracing xv6's system call no. 1 (fork) ~ no. 23 (proj2call)
- **Task #3:** Implement the user defined shell
  - Making the shell to execute commands and support > redirection

## ■ Submission Due

- Due: 5/21, Sunday 23:59
- No late submission is allowed

## ■ What to Submit (single tarball .tgz or .tar)

- System call files: syscall.c, syscall.h, sysproc.c, user.h, usys.S, Makefile
- User defined shell file: sh.c

## ■ Grading

- Total: 100 pts